

# Scheduling Algorithms

Instructor: William Zheng

Email:

[william.zheng@digipen.edu](mailto:william.zheng@digipen.edu)

PHONE EXT: 1745

# Lecture Outline


- Scheduling Intro
  - Preemptive and Non-preemptive algos
  - Types of Processes
  - Scheduling rubrics
- Non-preemptive Scheduling Algos
  - FCFS (First Come First Serve)
  - SJF
- Preemptive Scheduling Algos
  - Round-Robin
  - Multilevel Queues
  - Multilevel Queues with feedback

# Preemptive and Non-preemptive Scheduling




- Non-preemptive schedules *only* when
  - A process terminates
  - A process goes into waiting state
- Preemptive schedules *also* when
  - A hardware interrupt arrives
    - Process goes from waiting to ready state.
    - Process goes from running to ready state.
  - A software-generated interrupt (system call)
    - I/O request occurs requesting to access a file on hard disk

# Types of processes

- What processes/threads do?
  - Compute
  - I/O
- Types of processes/threads
  - I/O bound
  - Compute/CPU bound



# Scheduling Effectiveness

- **Waiting time**  

- **Turnaround time**  

- **Response time**  


**Turnaround time** – amount of time to execute a particular process





**Waiting time** – amount of time a process has been waiting in the ready queue

**Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

# First Come First Serve


- Schedule according to arrival order
- Implemented using a queue data struct
  - What's a queue data struct?
  - Basically a linked-list
    - Always read from front of linked list
    - Always add to end of the linked list

# FCFS waiting time example

- $P_1$  19 ms 
- $P_2$  7 ms 
- $P_3$  5 ms 
- Schedule 
- Wait time( $P_1$ ) = 0 ms
- Wait time( $P_2$ ) = 19 ms
- Wait time( $P_3$ ) = 26 ms
- Average wait time =
  - $(0 \text{ ms} + 19 \text{ ms} + 26 \text{ ms})/3 = 15 \text{ ms}$

# FCFS scheduling with I/O

- ▶  $P_1$  (10ms CPU, 4ms I/O, 3ms CPU)
- ▶  $P_2$  (3ms CPU, 4ms I/O, 3ms CPU, 4ms I/O, 2ms CPU)
- ▶  $P_3$  (5ms CPU, 5ms I/O, 8ms CPU)
- ▶ Schedule:









# Priority Scheduling

- A **priority queue** is a type queue with two operations
  - Put – place an item on the queue
  - Get – get the item with *highest* value item from the queue
- Same as FCFS scheduling algorithm, except that the queue is replaced by a priority queue
- What priority?
  - Based on feature of the jobs OR
  - A user-assigned priority value OR
  - Dynamic run-time priorities (pre-emptive versions)

# Shortest Job First

- Also known as *Shortest Remaining Time First*.
  - Shortest-next-CPU-burst algo
  - Associates with each process the length of the process's next CPU burst, rather than total length
  - Assigned to process with the smallest next CPU burst when CPU available
    - FCFS scheduling used to break the tie if the next CPU bursts of 2 processes are the same
- Non-preemptive priority scheduling
  - Always schedule the job with shortest job or shortest remaining time.
  - Minimizes average waiting time.

# SJF waiting time example

- ▶  $P_1$  19 ms 
- ▶  $P_2$  7 ms 
- ▶  $P_3$  5 ms 
- ▶ Schedule 
- ▶ Wait time( $P_1$ ) = 12 ms
- ▶ Wait time( $P_2$ ) = 5 ms
- ▶ Wait time( $P_3$ ) = 0 ms
- ▶ Average wait time =  
 $(12 \text{ ms} + 5 \text{ ms} + 0 \text{ ms})/3 = 17/3 \text{ ms}$

# SJF scheduling with I/O

- ▶  $P_1$  (10ms CPU, 4ms I/O, 3ms CPU)
- ▶  $P_2$  (3ms CPU, 4ms I/O, 3ms CPU, 4ms I/O, 2ms CPU)
- ▶  $P_3$  (5ms CPU, 5ms I/O, 8ms CPU)
- ▶ Schedule (according to next CPU burst):

Process in Ready Q before scheduling	$P_1$ 10ms $P_2$ 3ms $P_3$ 5ms	$P_1$ 10ms $P_3$ 5ms	$P_1$ 10ms $P_2$ 3ms	$P_1$ 10ms	$P_3$ 8ms $P_2$ 2ms	$P_3$ 8ms	$P_1$ 3ms
Scheduled Process	$P_2$ 3ms	$P_3$ 5ms	$P_2$ 3ms	$P_1$ 10ms	$P_2$ 2ms	$P_3$ 8ms	$P_1$ 3ms
Process in Waiting State after scheduling		$P_2$ 4ms	$P_3$ 3ms(of 5ms)	$P_2$ 4ms $P_3$ 2ms	$P_1$ 2ms	$P_1$ 2ms	

# Priority scheduling problem

- **Starvation:** a low priority job may sit on the queue indefinitely
  - Happens if there is always an incoming job request with a higher priority than the lowest priority job on the queue
- **Aging** as a solution:
  - if a job has waited on the queue for a sufficiently long time, its priority is increased by a set amount

# Pre-emptive scheduling


- Each job is allowed a time slice, or **time quantum**, in which to execute
- Once a time quantum has passed, or if the job has finished within that time, a different job gets a time slice
- A scheduling algorithm is used to determine the order in which jobs get a time slice.
  - **Round robin (RR) scheduling**
  - **Multilevel queue scheduling**
  - **Multilevel feedback queue scheduling**

# Round robin scheduling

- Uses a FIFO queue to process jobs in a first-come first-serve order
  - Get a job from the front of the queue
  - Let it execute for a maximum of one time quantum
  - If it does not complete within that time
    - Pause the job, and put the job at the back of the queue
  - Repeat

# RR waiting time example (5ms slices)

- ▶  $P_1$  7 ms
- ▶  $P_2$  19 ms
- ▶  $P_3$  5 ms
- ▶ Schedule



Waiting time: Sum of the periods spent waiting in the ready queue

1<sup>st</sup> part of waiting time for  $P_2$  is 5ms  
2<sup>nd</sup> part of waiting time for  $P_2$  is  $t_4 - t_2 = 7\text{ms}$ , total =  $5 + 7 = 12\text{ms}$ , i.e. total time spent in the ready queue for the remaining job.

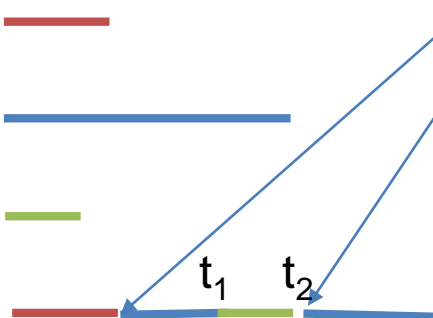
1<sup>st</sup> part of waiting time for  $P_1$  is 0ms  
2<sup>nd</sup> part of waiting time for  $P_1$  is  $t_3 - t_1 = 5(P_2) + 5(P_3) = 10\text{ms}$

- ▶ Wait time( $P_1$ ) = 10ms
- ▶ Wait time( $P_2$ ) = 12ms
- ▶ Wait time( $P_3$ ) = 10ms
- ▶ Average wait time =

$$(10\text{ ms} + 12\text{ ms} + 10\text{ ms})/3 = 32/3\text{ ms}$$



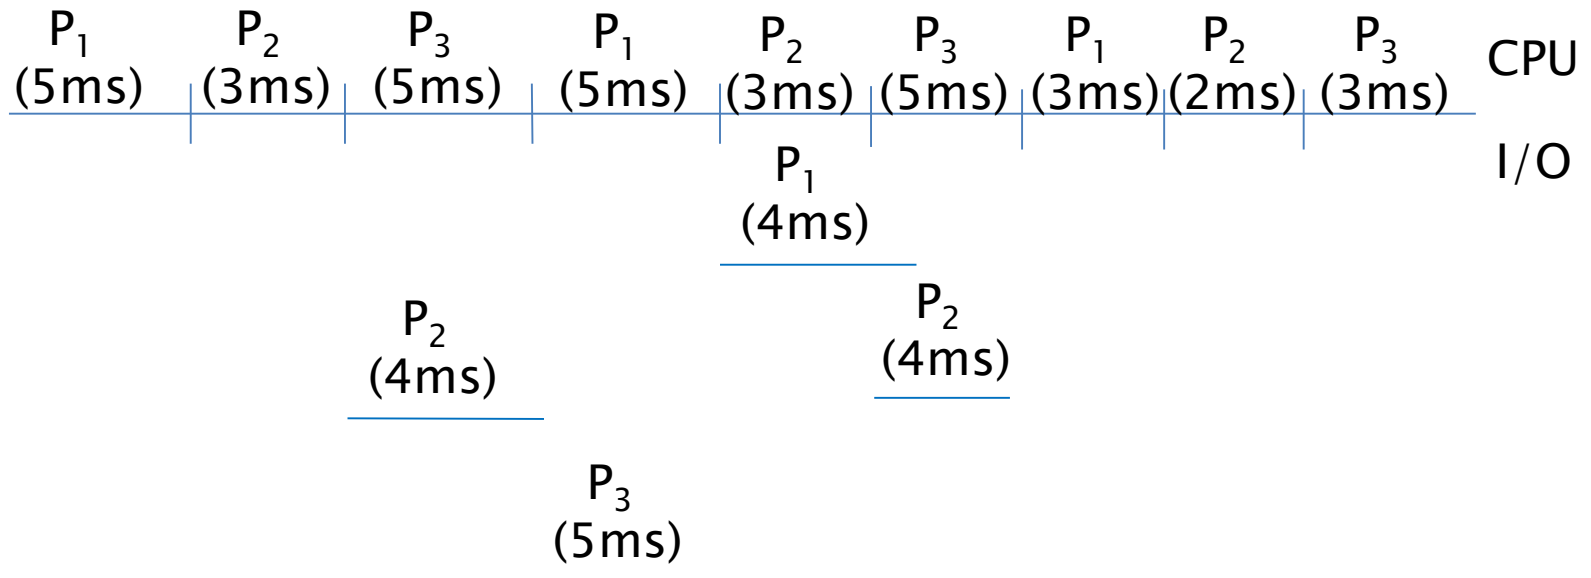
# RR waiting time example (10ms slices)

- ▶  $P_1$  7 ms —
- ▶  $P_2$  19 ms —
- ▶  $P_3$  5 ms —
- ▶ Schedule 
- ▶ Wait time( $P_1$ ) = 0ms
- ▶ Wait time( $P_2$ ) = 12ms
- ▶ Wait time( $P_3$ ) = 17ms
- ▶ Average wait time =  
$$(0 \text{ ms} + 12 \text{ ms} + 17 \text{ ms})/3 = 29/3 \text{ ms}$$

1<sup>st</sup> part of waiting time for  $P_2$  is 7ms  
2<sup>nd</sup> part of waiting time for  $P_2$  is  
 $t_2 - t_1 = 5(P_3) = 5\text{ms}$

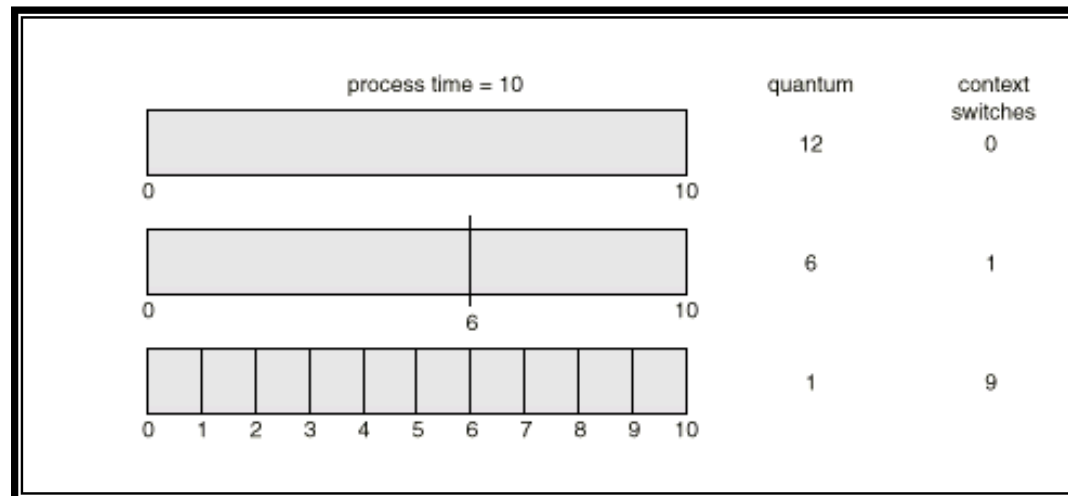
# RR scheduling with I/O (5 ms slices)

- ▶  $P_1$  (10ms CPU, 4ms I/O, 3ms CPU)
- ▶  $P_2$  (3ms CPU, 4ms I/O, 3ms CPU, 4ms I/O, 2ms CPU)
- ▶  $P_3$  (5ms CPU, 5ms I/O, 8ms CPU)
- ▶ Schedule:



# RR considerations

- The average waiting time decreases (in general) with decreasing time quantum size
- Context switching causes more processing time
- A balance must be made between time quantum size and context switching time
  - Want the time quantum to be “large” with respect to the context switching time



# Multilevel queue scheduling

- Ready queue is partitioned into separate queues
  - Jobs are put on different queues according to prioritization criteria
    - System jobs might have top priority
    - User jobs might have less priority
    - Background jobs might have least priority
- Each queue has its own scheduling algorithm
- Scheduling must be done between the queues. E.g.,
  - **Fixed priority scheduling**; (i.e., serve all from foreground then from background).  
Problem: **starvation**.
  - **Time slice** – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR, 20% to background in FCFS  
Problem: **how to split?**

# Multilevel feedback queue scheduling

- Uses multiple queues, as in multilevel scheduling
- However, a job may be moved between each of the different level queues, depending on CPU usage
- If a job has long CPU burst cycles, it is put on a lower priority queue

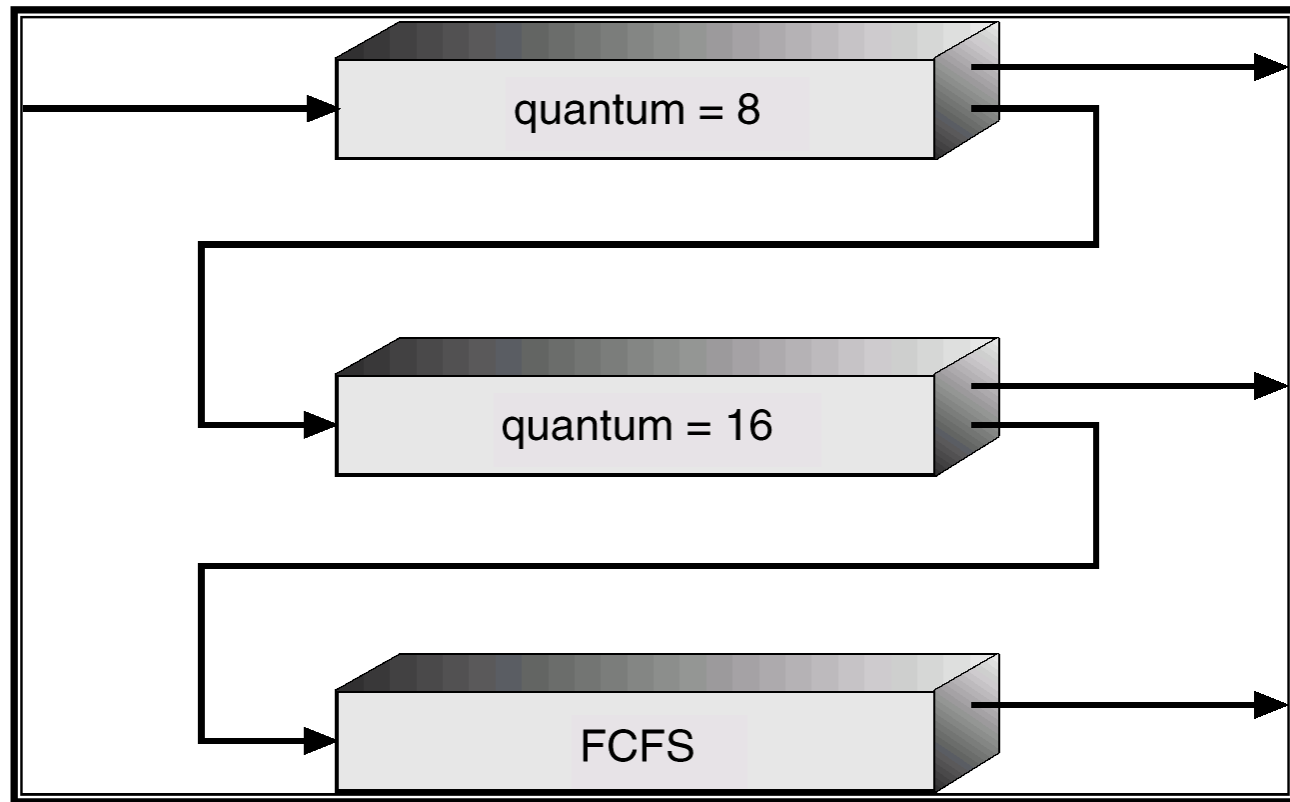
# Multilevel feedback queue scheduling

- Process can **move** between the various queues; aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
  - number of queues
  - scheduling algorithms for each queue
  - method used to determine when to upgrade a process
  - method used to determine when to demote a process
  - method used to determine which queue a process will enter when that process needs service

# Example of Multilevel Feedback Queue

- Three queues:
  - $Q_0$  – time quantum 8 milliseconds
  - $Q_1$  – time quantum 16 milliseconds
  - $Q_2$  – FCFS
- Scheduling
  - A new job enters queue  $Q_0$  which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue  $Q_1$ .
  - At  $Q_1$  job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue  $Q_2$ .

# Multilevel Feedback Queues






# Evaluation (1)

- Consider 5 processes arriving at time 0:

Process	Burst Time
$P_1$	10
$P_2$	29
$P_3$	3
$P_4$	7
$P_5$	12



Preemptive Round Robin Scheduling (time quantum=10ms):

Waiting time for  $P_1 = 0$  ms

Waiting time for  $P_2 = (t_1 - t_0) + (t_5 - t_2) + (t_7 - t_6) = 10 + 20 + 2 = 32$  ms

Waiting time for  $P_3 = t_2 - t_0 = 20$  ms


Waiting time for  $P_4 = t_3 - t_0 = 23$  ms

Waiting time for  $P_5 = (t_4 - t_0) + (t_6 - t_5) = 40$  ms


Average waiting time =  $(0 + 32 + 20 + 23 + 40) / 5 = 23$  ms

# Evaluation (2)


□ FCFS is 28ms:



□ Non-preemptive SJF is 13ms:



□ RR is 23ms:



Process	Burst Time
$P_1$	10
$P_2$	29
$P_3$	3
$P_4$	7
$P_5$	12

# Evaluation (3)

- Average turnaround time:  $(24+27+30)/3 = 27\text{ms}$

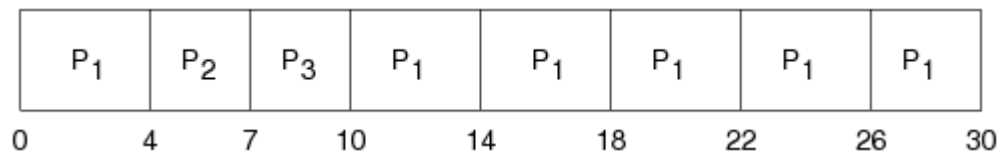
Process	Burst Time	Arrival	Start	Wait	Finish	TA
1	24	0	0	0	24	24
2	3	0	24	24	27	27
3	3	0	27	27	30	30



# Evaluation (4)

- average waiting time:  $(6+4+7)/3 = 5.67$
- average turnaround time:  $(30+7+10)/3 = 15.67$

Process	Burst Time	Arrival	Start	Wait	Finish	TA
1	24	0	0	6	30	30
2	3	0	4	4	7	7
3	3	0	7	7	10	10




# Turnaround Time and Response Time

- Turnaround time
  - the interval from the time of submission of a process to the time of completion
  - Sum of the periods spent waiting to get into memory, waiting in the ready queue, executing on CPU, and doing I/O
- Response Time
  - The time from submission of a request until the first response is produced
  - The time taken to start responding
  - Not the time taken to output the response

# Determining Length of Next CPU Burst

- Can only estimate the length – should be similar to the previous one
  - Then pick process with shortest predicted next CPU burst
- Can be done by using the length of previous CPU bursts, using exponential averaging
  1.  $t_n$  = actual length of  $n^{th}$  CPU burst
  2.  $\tau_{n+1}$  = predicted value for the next CPU burst
  3.  $\alpha, 0 \leq \alpha \leq 1$
  4. Define :  $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$ .
- Commonly,  $\alpha$  set to  $\frac{1}{2}$
- Preemptive version called **shortest-remaining-time-first**

# Prediction of the Length of the Next CPU Burst



CPU burst ( $t_i$ )	6	4	6	4	13	13	13	...
"guess" ( $\tau_i$ )	10	8	6	5	9	11	12	...